

Shameek Agarwal

✉ shameek.agarwal@gmail.com ☎ +916290885679 🌐 LinkedIn 🌐 Website

SUMMARY

Senior Software Engineer experienced in shipping production systems across data infrastructure, distributed systems, backend services, and frontend interfaces.

EXPERIENCE

New Relic, *Senior Data Platform Engineer* 07/2025 – Present

- Projected \$1M+ in annual savings on compute and licensing by leading the migration from DBT Cloud and Snowflake to DBT Core, Spark, Airflow and Iceberg. Conducted performance testing using TPCCH data to ensure Spark pipelines matched Snowflake throughput and latency.
- Reduced billing discrepancy Mean Time to Detect (MTTD) from months to ~1 hour by designing and shipping an automated, end-to-end billing validation framework adopted by 20+ teams.

New Relic, *Data Platform Engineer* 08/2024 – 06/2025

- Slashed pipeline delivery time from 2+ weeks to 2–3 hours by architecting a scalable ingestion framework with dynamic Airflow DAGs, Python, and Iceberg, onboarding 25+ analytical use cases.
- Deployed Airbyte into New Relic's infrastructure, migrating 10+ ingestion pipelines from Fivetran.
- Built a custom connector using Fivetran SDK for ingesting data from Billing Platform to Iceberg.
- Achieved 99.95% SLO and reduced PagerDuty alerts from 4/week to 1/month by adding freshness and volume tests in Python and New Relic, reducing incident detection from 2+ weeks to under 30 minutes.
- Ingested 65 GB/s of streaming data in real time by designing scalable Kafka Connect pipelines to fan out high throughput Kafka topics directly into Apache Iceberg, Amazon S3, and Snowflake.
- Deployed Spark on Kubernetes using KubeFlow and authored 10+ batch and streaming jobs.

Goldman Sachs, *Analyst (Full Stack Engineer)* 02/2022 – 07/2024

- Migrated 50GB+ of legacy records from PostgreSQL and enabled analytics via Snowflake by building a Spark based pipeline using the in house data platform.
- Automated external messaging by developing an email microservice using Spring Boot, Spring Security, and Hibernate, leveraging multithreading and in house object storage to handle attachments.
- Accelerated frontend build and release cycles by decoupling using micro frontends. Extracted shared components into unified libraries. Achieved using React, TypeScript and Webpack Module Federation.
- Established a single source of truth across 1M+ entities by consolidating company and investor records using Elasticsearch fuzzy matching, reducing duplicate profiles by 10%.

EDUCATION

Manipal Institute of Technology, *Bachelor's Degree in Computer Science* 2018 – 2022 | Udupi, India
9.64 CGPA

TECHNICAL SKILLS

Languages - Java, JavaScript, Python, Low Level Design, DSA, Object Oriented Design

Frameworks - Spring Boot, Spring Security, Hibernate, ReactJS

Databases - PostgreSQL, MySQL, Elasticsearch, MongoDB, Redis, SQL, NoSQL

Big Data - Spark, DBT, Iceberg, Snowflake, Hadoop, Kafka, Kafka Connect, Airflow, ETL, Warehousing

Cloud & Tools - Docker, Kubernetes, CI / CD, AWS, Terraform, IAC